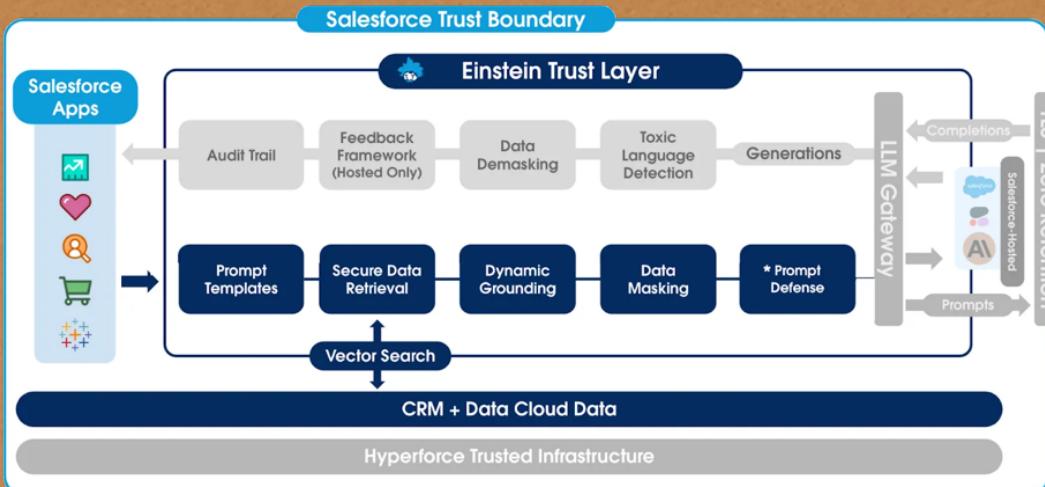


Agentforce

Humans with Agents drive customer success



Einstein Gen AI

It is a platform to securely connect your Salesforce data with the power of LLM to create relevant and customize responses.

Trusted Gen AI? → Salesforce Einstein Gen AI solutions are designed, developed and delivered based on below five principals

1. Accuracy
2. Safety
3. Transparency
4. Empowerment
5. Sustainability

WHY

"Review Gen AI Responses?"

1. Accuracy → Gen AI can sometime "hallucinate" and gives you results which are not grounded enough.
2. Bias and Toxicity → It can contain bias against historical collected data.

Note: If the response is not the expected output then you should regenerate it.

Einstein Requests are available in below environments -

- Enterprise, Performance and Unlimited edition with Einstein for Sales, Einstein for Service or Einstein for Platform add-on.
- All Einstein 1 Editions
- Marketing Cloud - Growth - Enterprise Edition and Marketing Cloud - Growth - Unlimited Edition

LLM Support → As of today Sep 2024 salesforce is having below listed LLM models to use with einstein

1. Anthropic Claude 3 Haiku
2. Azure OpenAI GPT-3.5 Turbo
3. _____ 16K
4. _____ 4 Turbo
5. OpenAI GPT-4
6. _____ 32K
7. _____ Turbo

Einstein Gen AI Limitations

→ In 1 minute you can have 70 LLM generations. This is a default rate limit by Salesforce and can be increased by talking to support.

→ If you don't see gen AI record field icon (**) that means prompt template is not supported for the field component.

High Level Org Setup

Step-1 Provision Data Cloud

Step-2 Turn on Einstein →
 |
 | Setup
 | | Quick Find
 | | Einstein Setup

Step-3 Configure the Trust Layer

Einstein Trust Layer

Einstein Trust Layer is a secure AI architecture built into Salesforce platform. It is a set of agreements, security technology and data privacy controls used to keep your company safe while you explore generative AI solutions.



Zero-Data Retention Policy

Your data isn't retained by third-party LLMs. We partner with Open AI and Azure Open AI to enforce the zero-data retention policy.

- No data is used for LLM model training or product improvements by third-party LLMs.
- No data is retained by the third-party LLMs.
- No human being at the third-party provider looks at data sent to their LLM.



Dynamic Grounding with Secure Data Retrieval

- Relevant information from a Salesforce record is merged with the prompt to provide context.
- Secure data retrieval of Salesforce data for grounding the prompt based on the permissions of the user executing the prompt.
- Secure data retrieval preserves in place all standard Salesforce role-based controls for user permissions and field-level security when merging grounding data from your CRM instance or Data Cloud.



Prompt Defense

- System policies help limit hallucinations and decrease the likelihood of unintended or harmful outputs by the LLM.
- System policies can vary for different generative AI features and use cases.



Data Masking

- Sensitive data is detected and masked before sending the prompt to the LLM.
- Data masking supports multiple regions and languages.
- You can select what must and must not be masked.



Toxicity Scoring

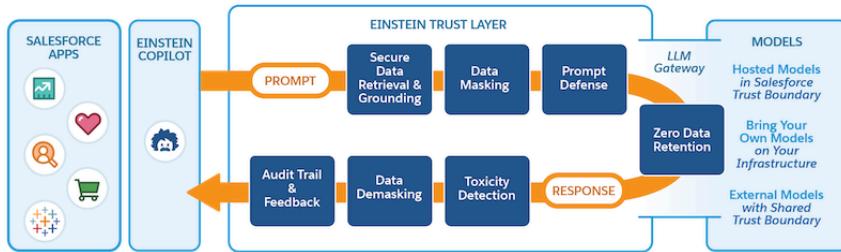
- Einstein Trust Layer scores content on toxicity.
- Toxicity scores are logged and stored in Data Cloud as part of the audit trail.



Audit ← This Audit Trail is in Data Cloud.

- Prompts, responses, and trust signals are logged and stored in Data Cloud.
- Feedback can be used for improving prompt templates.
- Provides pre-built reports and dashboards for analysis.

How Does It Work?



How to Setup Einstein Trust Layer?

1. Setup → Quick Find → Einstein Trust Layer
2. Enable Data Masking if it's turned off.
3. To change the default config you can click on "Configure Data Masking"

Notes: 1. By default Einstein Trust Layer identify and do the data masking for Personally Identifiable Information (PII) and Payment Card Industry (PCI) data in prompts.

2. Data Masking will replace the information with some placeholder text based on what it represents.
3. After the LLM returns the response the trust layer demask the data that was originally masked.
4. You can track the masked data by using the Audit Trail which is stored in Data Cloud.

Audit Trail → You can track your use of gen AI in your Salesforce org and can make sure that it complies with your privacy & security policies.

This needs to be turned on from Setup → Quick Find → Einstein Feedback.

A report package needs to be installed with it to access the audit trail report & dashboard data.

Data Masking Details



The following data has been masked. [Learn more in Salesforce Help.](#)

Placeholder Text	True Value
<ORGANIZATION_0>	Acme Corporation
<ORGANIZATION_1>	Global Media
<PERSON_0>	Howard Jones
<PERSON_1>	Carol White

OK

Prompt Builder

Prompt Builder helps you to create, modify, test and customize prompt templates.

Types of templates in Prompt Builder.

1. Sales Email Template → Draft personalised emails.
2. Field Generation Template → AI associated workflow to fields in lightning record page.
3. Flex Prompt Template → More flexible when it comes to generate prompts. It can accept upto 5 related objects data.

Navigate the Prompt Builder UI

Let's explore the Prompt Builder UI for a Sales Email prompt template.

The screenshot shows the Prompt Template Workspace for a Sales Email. At the top, there's a navigation bar with 'Prompt Builder', 'Draft an introduction Email', and 'Version 1'. Below it is a 'Resource' search bar with a placeholder 'Search for or select a resource to insert...' and a 'Contact' icon. To the right is a 'Configuration' panel with tabs for 'Template Properties' (selected) and 'Model Properties'. Under 'Template Properties', it shows 'Model Type: Standard' and 'Model: OpenAI GPT 3.5 Turbo'. Numbered callouts point to various parts: 1 points to the Resource search bar; 2 points to the Configuration panel; 3 points to the Preview section; 4 points to the Resolution section; 5 points to the Response section; 6 points to the 'Activate', 'Save As', and 'New' buttons in the Configuration panel; 7 points to the 'Resolution' section; and 8 points to the 'Response' section containing email preview text.

PROMPT TEMPLATE TYPE	WHAT IT CREATES
Field Generation	Generate content for record fields in Lightning Experience. In Salesforce, users click a button to run this prompt and populate the field with output. See Create a Field Generation Prompt Template and Field Generation Prompt Templates in Action .
Flex	Generate content for any business purposes that other templates don't cover. Flex prompt templates let you define your own resources. See Create a Flex Prompt Template and Flex Prompt Templates in Action .
Sales Emails	Generate a personalized customer email based on record data. See Create a Sales Email Prompt Template and Draft a Sales Email with Einstein Generative AI .
Record Summary	Summarize record data for a comprehensive view of a record. Record summary prompt templates are used by the Summarize Record standard invocable action. In Einstein Copilot, they're used by the Summarize Record standard copilot action. See Copilot Action: Summarize Record

User can use below ingredients to generate a prompt template.

1. Participants → Sender and Receivers.
2. Roles → Role of Sender and Receivers.
3. Relationship → Relationship between Sender and Receiver.
4. Instructions → Detailed instructions on what it must include and exclude while generating the template. Instruction must be with triple quote boundaries.
5. Guidelines → Proper guidelines to avoid hallucination.

Example: Generate a summary of open cases of the Astro's bakery Acc.

```
1 You're a support representative and you need to create a short summary of all open cases for the Astro's Bakery account.  
2  
3 When I ask you to summarize the open cases, you must strictly follow my instructions below.  
4  
5 Instructions:  
6 """"  
7 Summarize the open cases in one paragraph no longer than 500 characters. Mention how many open cases there are and what the case issues are.  
8 Use clear, concise, and straightforward language using the active voice and strictly avoiding the use of filler words and phrases and redundant language.  
9  
10 Use the following information to write the summary:  
11  
12 Case #1 details: The widgets we received are the wrong size. We needed a widget in size A, but we received a widget in size B.  
13  
14 Case #2 details: Our Widgets haven't been delivered. We can't track our order. According to the last tracking update, the Widgets should have arrived on June 5th. Now it's June 10th, and they still haven't arrived.  
15  
16 Do not attribute any positive or negative traits in the summary.  
17 """"  
18  
19 Now create the summary.
```

Result

```
21  
22 There are two open cases for the Astro's Bakery account. In Case #1, the issue is that the received widget is the wrong size. The requested size was A, but size B was received instead. In Case #2, the issue is that the widgets haven't been delivered and the order cannot be tracked. The expected delivery date was June 5th, but as of June 10th, the widgets have not arrived.  
23
```

Grounding Prompt Templates: Grounding provide LLM the context data to personalized the template.

We can use below approaches to ground the data.

- 1- Record Merge Fields
- 2- Flow Merge Fields
- 3- APEX _____
- 4- DMO _____
- 5- Related list _____
- 6- Record Snapshot
- 7- Retriever Augmented Generation (RAG) in Data Cloud.

Best Practices for Building Prompt Templates

Learn the best practices for building effective prompt templates.

REQUIRED EDITIONS

Available in: Lightning Experience

Available in: Enterprise, Performance, and Unlimited Editions with the Einstein for Sales, Einstein for Platform, or Einstein for Service add-on

- Make sure that your prompt templates are concise and easy to understand. Avoid using industry jargon or technical terms. Instead, use natural language and [conversation design guidelines](#).
- To give the LLM more contextual information, ask it to role play as a character, such as a sales or support representative. Then define the character's goal. For example, include language such as, "You are a marketing executive who wants to invite major customers to a live event."
- Iterate on your prompt templates. Try achieving the same goal using different templates to see how the parts impact the model's response. Get end-user feedback to see how well your prompt templates generate the desired response.
- Choose a style, and stick to it. When you use a consistent writing style in your prompt templates, the LLM generates consistent responses. Your writing style is shaped by your word choice, intensifiers, emojis, and punctuation.
- To help the LLM differentiate between context and instructions, create an instructions section in your prompt template. On a separate line, enter Instructions:, then surround your instructions with triple quotes (""""").
- Include direct instructions for the LLM to only generate the expected type of content. For example, if you want the LLM to draft an email, add instructions such as, "Follow these instructions strictly to generate only the message to be sent to the customer." These instructions prevent the LLM from generating a response about the process of creating content, instead of just generating the content that you want.
- Start with one of the templates in the [Example Prompt Template Library](#), and customize it to fit your specific needs. Study the language that the templates use, especially the text related to writing style. You can use similar phrasing in your own templates.

Numerical Limits

Prompt templates have these numerical limits.

Numerical Limits

LIMIT DESCRIPTION	LIMIT
Maximum template size	128,000 characters
Maximum number of template versions	50
Maximum number of merge fields	50
Maximum number of flow merge fields	5
Maximum number of Apex merge fields	5
Maximum number of related list merge fields	5
Maximum number of inputs in a Flex template	5

Audit Trail Isn't Supported

Creating or updating a prompt template isn't tracked in the Setup Audit Trail.

New Custom Objects and Custom Fields Aren't Immediately Available After Creation

When you create a custom object or custom field, it isn't available for use in Prompt Builder immediately. To use a new custom object or field immediately, log out and log back in.

Records Without a Name Don't Display in the Record Selector

Prompt templates with encrypted or missing name fields can't be previewed. Their records are unavailable in the test record selector.

Supported Object Types

Salesforce supports many object or field suffixes to represent different types of data. Prompt Builder supports only these types of objects and fields as resources.

- UI API-supported objects. For a full list, see [All Supported Objects](#).

 **IMPORTANT** The KnowledgeArticleVersion object is supported only for Flex Template types.

- Standard objects (no suffix)
- Standard fields (no suffix)
- Custom objects (_c)
- Custom fields (_c)
- Histories (_history)
- Related lists (_r)

Object Preview Limitations

You can't preview objects with no Name field or objects whose Name field is encrypted.

Limitations for Field Generation Templates

You can't use Field Generation prompt templates for rich text areas in knowledge entities (_kav). To populate a knowledge entity's rich text area with a Prompt Builder template, use a Flex prompt template, Flow, or Apex.

Limitations for Flex Templates

You can't import or export metadata from flex templates, including metadata related to flex templates. For example:

- Apex classes that reference flex templates
- Flows that reference flex templates
- Change sets to move flex-related artifacts between orgs

To delete flex templates that use flows or Apex classes, first remove any flows and Apex resources in all versions of the template. Second, delete any flows and Apex classes that reference the prompt template. Then delete the flex template.

Limitations for Einstein Search

If a prompt uses an Einstein Search retriever, the change sets don't include the retriever or search index metadata. You must manually create the retriever in the destination org before deploying the retriever. This rule applies to change sets and Metadata API deployments.

User Permissions

- Prompt Template Manager Permission Set → Create & Manage prompt templates in Prompt Builder,
- Prompt Template User Permission Set → To access and run prompt templates outside Prompt Builder .

Available Objects for Merge Fields as per Templates.

Template	Merge Fields Objects
Sales Email Template	Recipient (Contact or Lead) Sender (the User) Current Org (Sender's org)
Field Generation Template	Current User Current Org The Object that user add in config
Flex Prompt Template	Current User Current Org Objects that user associated with template config.

Open Merge Fields

1. The apex class must contain one method annotated with **InvocableMethod**
2. The InvocableMethod must have capability type details as per your selected template.

PROMPT TEMPLATE TYPE	CAPABILITYTYPE
Sales Email	PromptTemplateType://einstein_gpt_salesEmail
Field Generation	PromptTemplateType://einstein_gpt_fieldCompletion
Flex	FlexTemplate://template_API_Name
Record Prioritization	PromptTemplateType://einstein_gpt_recordPrioritization
Record Summary	PromptTemplateType://einstein_gpt_recordSummary

3. The method input parameter must have a **List<Request>** class defines an **InvocableVariable** member for each input defined by the prompt template type.
4. The method output has a **List<Response>** type containing a single string member named **Prompt**.

Related List Merge Fields

Considerations, Guidelines, and Limitations

Consider these points when grounding your prompt templates with related lists in Prompt Builder.

- Activities related lists on objects such as Account and Case are not supported.
- If you configure a related list in a prompt template and a user doesn't have the related list on their parent object's page layout or Read permission on the associated object, no related list data is sent to the LLM.
- Only related lists available on the object's page layout are available for grounding in Prompt Builder. If an object has custom record types, only related lists on the master record type's layout are available.
- The correct syntax when manually typing (as opposed to using the resource picker) an expression for, as an example, a Contacts related list on the Account object is **{!\$RelatedList:Account.Contacts.Records}**. After typing this expression, Prompt Builder simplifies the UI presentation, leaving only Related List: Account.Contacts.
- In situations where a related list yields zero results, the prompt response still responds successfully but without adding any related list content to the response.
- The User and Organization related lists aren't supported, due to a limitation involving global variables.

Ground with Record Snapshot

Record Snapshot brings a new way to ground your prompts by automatically including data available on the User's page layout.

Considerations and Limitations

Consider these points when grounding your prompts with Record Snapshots.

- Dynamic Forms aren't used in Record Snapshots. However, if Dynamic Forms are activated for an object used by Record Snapshots, the underlying page layout is still used.
- Prompt Builder restricts the quantity of data received from each source to avoid overflowing the prompt token limits.
- "Empty" data (fields without values, sections without limits, and so on) are filtered out.
- Prompt Builder imposes these limitations on each data source used with Record Snapshots.

DATA SOURCE	LIMIT	CRITERIA
Page layout	None for the number of fields returned	
Compact page layout	None for the number of fields returned	
Related Lists	The first 6 related lists from the page layout are used, and up to 5 records are used per related list.	
Tasks	The next 5 upcoming tasks are used.	Tasks are selected based on ActivityDate (ascending order), Priority (descending order), and ModifiedDate (descending order).
Events	The next 5 upcoming events are used.	Events are selected based on start date (ascending order).
Emails	The 5 most recent emails are used.	

Ground with Retriever Augmented Generation (RAG) in Data Cloud

Using RAG you can use large subset of data like Knowledge Articles, Emails, Chat transcript.

After adding an active retriever to a prompt template, you can configure the dataset to fine-tune the data returned. In Prompt Builder, define search parameters in the Configuration side panel to adjust the retrieved dataset. These parameters include:

- Search Text: A fully dynamic field that uses merge fields from your prompt template inputs or free text to query the retrieved data. The search text input to the retriever is limited to globals and prompt inputs. To insert merge fields, use the resource picker above the search text area.
- Output Fields: The returned fields from the retriever that are included in the search results. The output fields vary and are determined during the retriever configuration in Einstein Studio. By default, all fields defined during the retriever configuration are included in the resolution.
- Number of Results: The maximum number of results that are returned.

If no value is specified for the list of output fields or the number of results, Prompt Builder uses the values set when the Einstein Search Retriever was created. Access to retrievers and their data is controlled by the [Data Cloud permission sets](#).

Data returned from the retriever is shown in the Prompt Builder Preview Resolution in JSON format. The prompt template response is then customized based on the search criteria.

Considerations

Consider these points when grounding your prompt templates with retrievers in Prompt Builder.

- The search text field for retrievers is limited to [globals](#) and [prompt inputs](#). It cannot use other sources, such as related list, Flow, and Apex.
- Retrievers that use a placeholder aren't available within Prompt Builder.
- If two Einstein Search merge fields are identical, their configurations are the same. They can't be configured independently.

Notes → * Flex Prompt Template Supports adding [FreeText](#) input box. This functionality allow you to handle scenarios where you want your templates to receive free text.

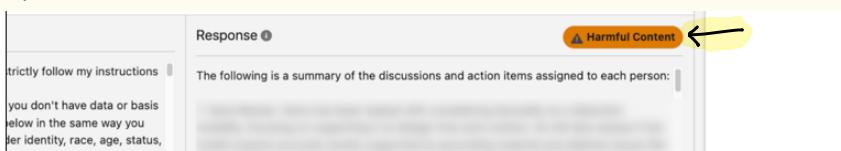
On the UI or Backend where you are populating values in other reference fields, you can add values in the free text area as well.

* While previewing prompt templates each time you click [Preview](#) button, the LLM is going to generate a new response. Even if you haven't changed your prompt template.

To assess a response's effectiveness, it's helpful to ask yourself questions about these topics.

TOPIC	QUESTIONS
Relevance	Is the response relevant in this context? Does the response fit in with the conversation or content that surrounds it?
Goal Completion	Does the response fulfill the goals of the prompt? Does it address everything that the prompt requests?
Style and Tone	Is the style, voice, and tone of the response appropriate? Is the response's vocabulary and punctuation correct?
Factual Accuracy	Does the response correctly incorporate the grounded data? Is the information in the response complete and accurate? Does the response contain redundant, excess, or erroneous information?
Consistency	How varied is the response? When you regenerate the response without changing the prompt template, how does the response change? How does the response change when you ground the prompt with different data?
Toxicity	Is the response safe? Does it avoid potentially harmful content, such as offensive, disrespectful, or abusive language? LLMs are trained on huge amounts of data, which puts the model at risk for producing toxic verbiage that leaks into your responses.
Bias	Does the response reflect fairness and inclusivity? Does it assume the gender identity of a person based on their name alone, sideline participants with disabilities, or displace assumptions about race or socioeconomic status? LLMs are trained on huge amounts of data, which puts the model at risk for producing biased verbiage that leaks into your responses.

* If the response contains toxic content like text with violent, abusive, hateful or sexual language, imagery or theme that can harm individuals or groups, The response is going to show you the harmful content badge.



If the Harmful Content badge appears, update and retest the prompt template to remove the toxic content. If the reassessed template no longer contains toxic content, the badge is removed.

* While working with larger prompt text the Prompt Builder will give you a summary generated automatically in the resolution panel. You can keep using the summarized prompt or create a similar prompt.

* We can use Example Prompt Template Library as a starting point to create our prompt templates.

Prompt Builder supports these languages for generative responses.

LANGUAGE	LOCALE	CODE
Dutch	Netherlands	nl_NL
English	United States	en_US
French	France	fr_FR
German	German	de_DE
Italian	Italy	it_IT
Japanese	Japan	ja_JA
Portuguese	Brazil	pt_BR
Spanish	Spain	es_ES
	Mexico	es_MX
Swedish	Sweden	sv_SE

Prompt Builder instructs the LLM to optimize responses for Spanish (Mexico) and Portuguese (Brazil) locales. Some models generate responses that are close variants but not exact matches for the requested locale.

* A prompt template response is automatically rendered into the user's language (if supported) or English, even if the prompt template itself is in another language, or if it instructs the LLM to respond in a specific language.

Agentforce : Agents and Copilot

Einstein copilot is a bot/Agent which takes questions posted by users in natural language and provides relevant answers drawn from secure, proprietary company data.

In Salesforce, your copilot is trusted, natural language, and conversational.

- Trusted: A copilot can take action in your Salesforce org based on a user's request. The copilot respects standard Salesforce access controls such as licenses and permissions, so it always acts securely. Einstein Copilot includes AI guardrails designed to help meet ethical standards, and it's integrated with the Einstein Trust Layer, a secure AI technology natively built into the Salesforce Platform.
- Natural Language: Instead of clicking around the Salesforce UI, users can express their questions or instructions as if they're talking to a human. The copilot, powered by a large language model (LLM), interprets the request and responds to the user in natural language.
- Conversational: A copilot can have multiple back-and-forth interactions with a user. Each user request or instruction is understood in the context of an ongoing conversation, so a user can ask follow-up questions or make related requests.

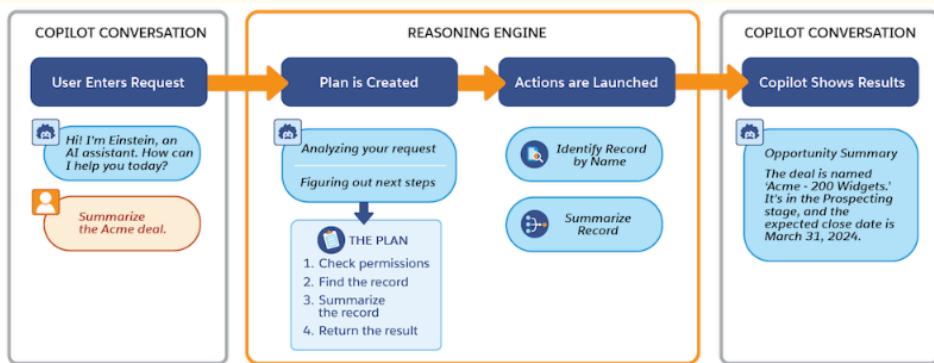
Setting up a copilot doesn't require any special expertise or coding skills. Just enable Einstein Copilot, and your copilot can assist your users with common tasks in Salesforce, such as summarizing records, drafting emails, and finding records.

It's also easy to extend your copilot by using preexisting Salesforce platform functionality. For example, if you have a flow in Salesforce that can make product recommendations, you can add that capability to your copilot as a custom action with a few clicks.

* Available in Enterprise, Performance and Unlimited editions with Einstein for Sales, Einstein for Service, or Einstein Platform add-on.

Einstein Copilot has five components

Copilot Topics Actions The Reasoning Engine LLM



How it works?

1. Identify the user's intent
2. Generate a Plan
3. Executing the Plan
4. Generating the response
5. Keeping the conversation going.

Note: Einstein Copilot respects standard Salesforce access controls such as licenses, permissions, FLS and Sharing Settings.

Einstein Copilot Analytics → It is a Dashboard which provide insights on the usability of your copilot.

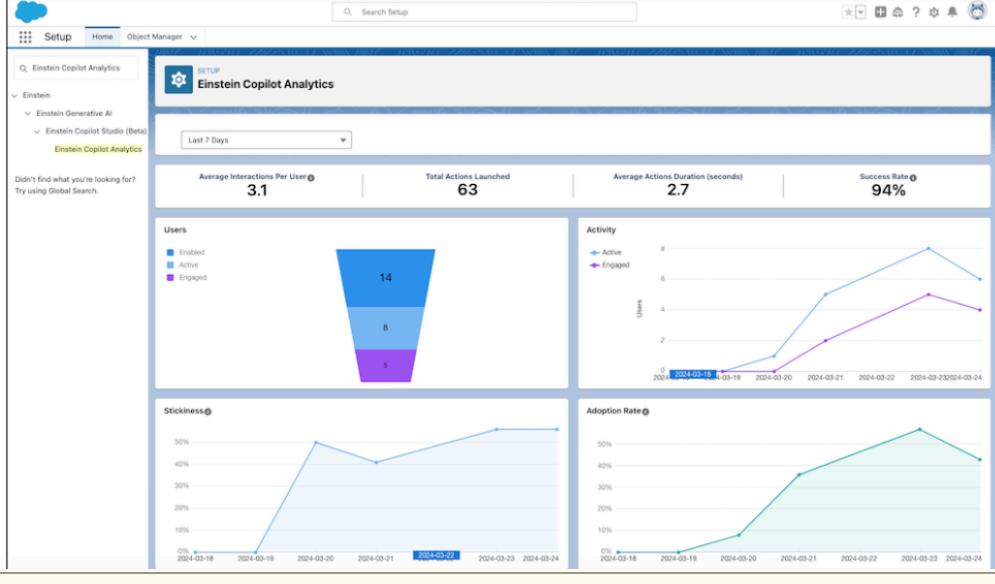
The information gathered and shown on the analytics dashboard is broken down into user data and action data. User data is aggregated and includes:

- Enabled users
- Active users
- Engaged users
- Repeat users
- Average interactions per user

Actions-related data includes both aggregated and individual action information. The dashboard includes:

- Total actions launched
- Average action duration
- Overall action success rate
- Specific action duration and success rate

To see the dashboards, from Setup, in the Quick Find box, enter Agent Analytics, and then select Agent Analytics.



Limitations

Einstein Copilot Support

- The Einstein panel is available in Lightning Experience, the Salesforce mobile app, the Field Service mobile app, and Sales Cloud Everywhere.
- Einstein Copilot supports English only. Some actions and associated prompt templates can be triggered with user messages in languages other than English and can return outputs in those languages, but this functionality isn't supported.
- Einstein Copilot supports OpenAI GPT-4o for planner service calls. Copilot actions can make calls to other predefined LLMs. Bringing your own model isn't supported, but custom actions that execute prompt templates can use any LLM supported in Prompt Builder.
- Einstein Copilot supports [custom and standard User Interface API-supported objects](#). Object support for individual standard copilot actions can vary. See the reference documentation for an action for more information.
- Einstein Copilot doesn't support the thumbs up/thumbs down component to collect and store user feedback.
- Einstein Copilot has limited API support.
- Einstein Copilot isn't available for Government Cloud.

Einstein Copilot Limits

- For best performance, we recommend assigning [no more than 15 actions to a topic](#).
- You can have up to [one copilot](#) (Einstein Copilot for Salesforce). Versioning isn't supported.
- Conversation history applies to the current session only. Messages from up to the most recent six turns (including copilot and user messages) are used as context. To keep the context focused on CRM use cases, messages such as "hello" aren't included.

Considerations for Einstein Copilot Conversations

- Einstein Copilot is optimized for questions or requests that are related to a specific action, rather than open-ended questions that are open to interpretation. Copilot supports [Salesforce-provided standard actions](#) and custom actions.
- Einstein Copilot sessions are specific to a browser tab. A new session is started when a user refreshes their browser tab or opens Einstein Copilot in another tab.
- Deactivating a copilot interrupts any ongoing user conversations. Users aren't notified that a copilot is deactivated. The copilot sends the system error message as a response to any user messages. The Einstein panel is available until the user closes it, and it can't be reopened until the copilot is active.
- For any user request that returns a list, a copilot shows up to 50 records when the user selects [View More](#), even if the number of records is larger than 50. The 50 records are shown in a record page, not in the copilot conversation.
- Page context is supported for some standard copilot actions. If a supported standard action is triggered by a user's request, Einstein Copilot is aware of which page the user is on. Page context is not supported for custom copilot actions.
- In general, Einstein Copilot supports one intent per utterance, so users can't make requests that span multiple intents. Users can break more complex requests into individual messages. However, the Draft or Revise Email action can combine actions to address multiple user intents.
- Copilot conversations aren't personalized based on past sessions. See Einstein Copilot Limits.
- Einstein Copilot can't make proactive suggestions. Users must enter a request or ask a question to start a conversation with a copilot.
- Einstein Copilot supports small talk, but topics and number of turns are limited to ensure trust and prioritize enterprise use cases. Great for CRM! Awkward at parties.



NOTE sObject input types are currently only partially supported due to inconsistent behavior. A fix is on the way. In the meantime, we recommend using a flow to wrap a prompt template and creating a custom action from the flow. This scenario requires detailed action instructions, so test and iterate on your instructions carefully.

- Outputs: text type outputs (for example, a two-sentence summary of the opportunity)
- Einstein Copilot custom actions support only [custom](#) and [standard](#) User Interface API-supported objects.
- Custom actions that reference an Apex class or flow support only [primitive data types](#). Collections aren't supported.

How to Setup, Manage User Access and Configure Copilot's settings?

Setup:

USER PERMISSIONS NEEDED

To enable Einstein Copilot for Salesforce:

Manage Einstein Copilot for Salesforce
OR
Customize Application
OR
Modify Metadata

Before you begin, turn on Einstein Generative AI.

1. From Setup, in the Quick Find box, enter Agent, and then select **Agents**.

If you don't see **Agents** in Setup, verify that Einstein Generative AI is enabled for your org. If that's not the issue, make sure you have the right license and user permissions for Einstein Copilot.

2. Turn on **Einstein Copilot for Salesforce**.

When you're ready, give users access to Einstein Copilot for Salesforce.

To enable permissions for a user:

1. From Setup, in the Quick Find box, enter **Users**, and then select **Users**.
2. In the list view, click the user's name to view the user details page.
3. Scroll down to the Permission Set Group Assignments section, and then click **Edit Assignment**.
4. Select the **Einstein Copilot for Salesforce User** permission set group, and then click **Add**.
5. Click **Save**.

If you prefer using permission sets to manage user permissions, assign the Einstein Salesforce for Copilot User and Prompt Template User permission sets.

Your copilot isn't visible to users until you activate it.

Configurations

Manage Copilot Settings

Use the copilot detail page and the Agent Builder to view and define important copilot settings such as event logs, system messages, and conversational style.

REQUIRED EDITIONS

Available in: Lightning Experience

Available in: Enterprise, Performance, and Unlimited Editions with the Einstein for Sales, Einstein for Service, or Einstein Platform add-ons.

- [Enable Enhanced Event Logs](#)

Event logs capture the events in a copilot conversation to help you test and troubleshoot your copilot. Enable enhanced event logs to store conversation data and view all conversation activity in one place.

- [Define System Messages](#)

Create welcome and error system messages that are sent when your users begin using copilot and if they encounter system errors. When you customize your welcome message, we recommend introducing the copilot as a bot or AI assistant, so your users know they aren't chatting with a human.

- [Update Copilot Language Settings](#)

Change the tone of your copilot's messages to better suit your company's brand. Einstein Copilot is available in English only.

Note: In the Copilot language settings you have three tones

1. Formal: Polite and Professional

2. Neutral: Unbiased and Nonevaluative

3. Casual: Relaxed and Informal

Customize Copilot with Topic and Actions.

Topic → Copilot Topics are basically a category of actions related to a particular job to be done.

Below are the ingredients of a topic

Parts of a Topic

FIELD	DEFINITION	EXAMPLES
Classification Description	1–3 sentences that describe what a topic does and the types of user requests that should be classified into this topic. A copilot uses this description to determine when to use a topic in a conversation. When a user enters a message, the copilot identifies the user's intent, compares it to the topic name and classification description, and uses the best matching topic.	Answers questions about our company's business policies, procedures, or products. Interacts with the user to draft or revise content from Salesforce CMS in the content builder. Answers questions related to a user's order status, return status, or making a return.
Scope	A specific job description for a topic. After a copilot has selected a topic based on user input, the scope gives a copilot more information about what it's able to do with a topic.	Your job is to only handle tasks related to setting a closing strategy to close deals. This may include answering product-related pricing questions, briefing users on customer sentiment based on conversations, suggesting a close plan to help users close deals, or giving advice based on video and voice calls. Your job is to only educate the user with their reservation questions and only handle changing upcoming or existing reservations by resending the reservation confirmation email. You aren't able to make new reservations.
Instructions	Instructions help a copilot make decisions about how to use the actions in a topic for different use cases. Each instruction is a single topic-specific guideline. Usually phrased as "Always...", "Never...", "If x, then y...", "As a first step..." Instructions are non-deterministic, so they shouldn't include sensitive or deterministic business rules in topic instructions.	Always clarify what type of mobile device (iOS or Android) before using the Answer Questions with Knowledge function to retrieve troubleshooting information. Never ask the user for an ID. Instead, use the name or ask for the name of the record to use IdentifyRecordByName to get the ID to call other actions. If a user asks about the status of their order or order confirmation, ask if they have their order number or their email. If a user doesn't have their email or order number, ask if they'd like to escalate to an agent or create a case.

In addition to these fields, a topic contains actions that are relevant to the topic's use case. Each action in a topic is a task that your copilot can perform when that topic is selected in the conversation. For example, an Order Management topic can contain:

- Get Order by Order ID
- Get Order by Email Address
- Look Up Return Status
- Create Return Label

Best Practices for Writing Topic Instructions

- When writing topic instructions, don't include sensitive or deterministic business rules (for example, "Don't refund an order unless it was within 30 business days."). To ensure consistent performance, build these requirements into the functionality of the reference action itself (for example, the Apex class or flow that a copilot action references). Instructions are non-deterministic and don't replace the need for coded business rules.
- Provide context based on the task you want the instruction to handle. For example, for a topic that handles watch repairs, the instructions should specify that the repairs are specifically for watches and spare link parts.
- To help your copilot understand what your users are asking for, include examples or sample inputs and outputs.
- Instructions are sensitive to language choice, so aim to create instructions that are precise, detailed, and direct, and then carefully test and iterate on them. In general, a copilot tends to strictly follow strong language such as "always" and "never." On the other hand, a copilot can apply instructions inconsistently or incorrectly when language is ambiguous.
- Before adding a new instruction, review your existing instructions to ensure that there are no conflicts. Contradicting instructions can cause errors and degrade copilot performance.

Create a new Topic

USER PERMISSIONS NEEDED

To create a custom topic for Einstein Copilot:

Manage Einstein Copilot for Salesforce

OR

Customize Application

OR

Modify Metadata

Before you begin:

- Identify the actions you want to add to your topic and [create any custom actions](#).
- [Review the parts of a topic and best practices for writing topic instructions](#).

1. From the Agents Setup page, launch your copilot in the Agent Builder.
2. [If your copilot is active, deactivate it](#).
3. From the left rail, select **Topics**.
4. On the Topics panel, from the dropdown menu, select **New Topic**.
5. Enter your topic's details, including your topic's name, classification description, scope, and at least one instruction. Then click **Next**. To learn more about these fields, see [Einstein Copilot Topics](#).
6. Select the actions you want to add to your copilot, and then click **Finish**.

Your custom topic is assigned to your copilot. To test your custom topic, restart the preview conversation to apply your changes.

Before activating your copilot and making your topic available to your users, carefully test and iterate on your topic's details, especially your topic's instructions. To test your topic, open the copilot in the Agent Builder and start a preview conversation. Enter utterances that you expect to trigger your topic and actions, and then adjust your topic's details based on your results.

You can view or edit your custom topic from the Topics panel in the Agent Builder. To view or edit your topic's details, click the topic name and select the **Topic Configuration** tab. To view, add, or remove copilot actions from your topic, select the **This Topic's Actions** tab.

Add a Topic from the Asset Library

With standard topics, your copilot can start handling common business use cases quickly.

REQUIRED EDITIONS AND USER PERMISSIONS

Available in: Lightning Experience

Available in: Enterprise, Performance, and Unlimited Editions with the Einstein for Sales, Einstein for Service, or Einstein Platform add-ons.

USER PERMISSIONS NEEDED

To assign a standard topic to a copilot:

Manage Einstein Copilot for Salesforce

OR

Customize Application

OR

Modify Metadata

1. From the Agents Setup page, launch your copilot in the Agent Builder.
2. If your copilot is active, [deactivate it](#).
3. From the left rail, select **Topics**.
4. On the Topics panel, from the dropdown menu, select **Add from Asset Library**.
5. Select one or more standard topics you want to add to your copilot, and then click **Finish**.

The topics you've added are assigned to your copilot and visible on the Topics panel in the Agent Builder. To test your new topic, restart the preview conversation to apply your changes.

Actions → Actions are the steps/Processes that Copilot is going to follow to get the job done.

Standard Actions

Salesforce provides some actions out of the box, and these actions are called standard actions. The benefit of including standard actions by default is that your copilot is ready to help users with many common Salesforce tasks as soon as you enable Einstein Copilot.



NOTE Some Einstein Copilot actions are in beta and have limited functionality, as further described in the Documentation. Including them in a copilot is part of the Services and will consume Einstein Requests if enabled and used.

Some standard actions are critical for the basic functionality of the copilot. They're considered system actions. They're often used to retrieve the information required to execute another action—for example, the ID of a record when a user refers to it by name. You can't remove system actions from your copilot.

Many standard copilot actions are available to all users who have permission to access Einstein Copilot. Other standard actions are built to work with specific clouds or products, so they require another license or permission. On the Copilot Actions page and in the Copilot Actions Library, you see only the actions that are available for your Salesforce org.

Custom Actions

You can also create custom actions to give your copilot more abilities so it can assist with tasks specific to your business. Access to a copilot custom action depends on how the action is built.

Copilot Action Assignments

A copilot uses only the actions that are assigned to it. To assign an action to a copilot, an action must be added to a topic. An action can be added to multiple topics.

View all of the actions that are assigned to your copilot on the Actions tab of your copilot's record page or on the actions panel in the Agent Builder.

View and manage the actions that are assigned to a topic from the Topics panel in the Agent Builder.

- [Add an Action to a Topic](#)

To use an action with a copilot, add it to a topic. An action can be added to multiple topics.

- [Extend Your Copilot with Custom Actions](#)

Custom actions help you extend your copilot so it can assist users with tasks specific to your business.

- [Copilot Actions and Large Language Model Use](#)

The cost associated with a copilot action depends on the number and size of large language model (LLM) calls, which vary by action and use case.

Add an Action to a Topic

To use an action with a copilot, add it to a topic. An action can be added to multiple topics.

REQUIRED EDITIONS AND USER PERMISSIONS

Available in: Lightning Experience

Available in: Enterprise, Performance, and Unlimited Editions with the Einstein for Sales, Einstein for Service, or Einstein Platform add-ons.

USER PERMISSIONS NEEDED

To assign actions to a topic:

Manage Einstein Copilot for Salesforce

OR

Customize Application

OR

Modify Metadata

For best performance, we recommend assigning no more than 15 actions to a topic.

You can also add an action to your copilot when you [create a custom topic](#).

1. From the Agents Setup page, launch your copilot in the Agent Builder.
2. If your copilot is active, [deactivate it](#).
3. From the left rail, select **Topics**.
4. On the Topics panel, click the name of the topic you want to add an action to, and then select the **This Topic's Actions** tab.
5. Click .
6. Select the action or actions you want to add to your copilot, and then click **Finish**.

The actions you've added are assigned to your copilot and are visible on the Topics panel of the Agent Builder. To test your topic with your newly assigned actions, restart the preview conversation to apply your changes.

To remove a copilot action from your copilot, from the Topics panel of the Agent Builder, click the name of the topic you want to remove an action from, and then select the **This Topic's Actions** tab. From the dropdown menu next to the action, select **Remove from Topic**.

Extend Your Copilot with Custom Actions

Custom actions help you extend your copilot so it can assist users with tasks specific to your business.

REQUIRED EDITIONS

Available in: Lightning Experience

Available in: Enterprise, Performance, and Unlimited Editions with the Einstein for Sales, Einstein for Service, or Einstein Platform add-ons.

When you create a custom action, you build it on top of platform functionality you want to make available in Einstein Copilot, such as **invocable Apex classes, autolaunched flows, or prompt templates**. Adding custom actions lets you customize your copilot and get more mileage out of your current Salesforce platform capabilities.

Access to a custom copilot action depends on the type of Salesforce action it references. For example, if a custom action was built using a flow, the custom action adheres to the permissions, field-level security, and sharing settings configured in the flow.

- **Create a Custom Copilot Action**

With custom actions, you can extend your copilot with Salesforce functionality so it's ready to assist users with tasks specific to your business. Build custom actions using invocable Apex classes, autolaunched flows, and prompt templates.

- **Best Practices for Einstein Copilot Action Instructions**

Action instructions tell a copilot what your action does and when to use it in a conversation. When you create a custom action, carefully plan and test your instructions to ensure that your action performs as expected.

Create a Custom Copilot Action

With custom actions, you can extend your copilot with Salesforce functionality so it's ready to assist users with tasks specific to your business. Build custom actions using invocable Apex classes, autolaunched flows, and prompt templates.

REQUIRED EDITIONS AND USER PERMISSIONS

Available in: Lightning Experience

Available in: Enterprise, Performance, and Unlimited Editions with the Einstein for Sales, Einstein for Service, or Einstein Platform add-ons.

USER PERMISSIONS NEEDED

To create a custom action for Einstein Copilot:	Manage Einstein Copilot for Salesforce OR Customize Application OR Modify Metadata
---	--

To create or modify prompt templates:	Prompt Template Manager
---------------------------------------	-------------------------

Before you begin:

- Review the [considerations for custom actions](#).
- Create an invocable Apex class, autolaunched flow, or prompt template. You'll configure one of these actions to use with Einstein Copilot.

1. From Setup, in the Quick Find box, enter Agent, and then select **Agent Actions** to view all the actions in your Salesforce org.
2. Click **New Copilot Action**.
3. From the Reference Action Type dropdown, select the type of action that you want to use to build your copilot action, such as a flow or prompt template. Then select the action.
4. The copilot action label and API name are populated with the reference action name and API name. Review these fields and, if necessary, edit the values so they accurately describe the action. Then click **Next**.
5. The instructions for the custom action and each input and output are populated with the descriptions from the reference action. Review these fields and make changes. See [Best Practices for Einstein Copilot Action Instructions](#).
- We copy over the descriptions from the reference action to give you a head start, but copilot action instructions differ from traditional descriptions. A copilot uses a large language model (LLM) and the planner service to determine when to launch an action in a conversation. The instructions for the action, inputs, and outputs tell a copilot what your action does and when and how to use it. Effective instructions vary by action and use case. Plan to test and iterate on your instructions to ensure that your action performs as expected.
6. For each input and output, specify any applicable settings.

SETTING	DETAILS
Collect data from user	Require this parameter to be provided by the user.
Filter from copilot action	When you create a copilot action, all inputs and outputs from the reference action are added to the copilot action. Enable this setting when you don't want to use a parameter from the reference action with your copilot action. At least one output must be used with the copilot action.
Require input	Require this parameter in order to execute the action. All other inputs are treated as optional. If an input is required by the reference action, this setting is enabled by default and read-only.
Show in conversation	Allow a copilot to include this parameter in a response to a user. At least one output must be available to show in a copilot conversation.

 **TIP** When an action makes a change to a record, you can require a copilot to ask the user to confirm the change before the copilot can execute it. On the record page for the action, enable the **Require user confirmation** setting.

7. Click **Finish**.

After you create your action, [assign it to a topic](#) to add it to your copilot.

To test your action and preview how the output appears in a copilot conversation, open the copilot in the Agent Builder and start a preview conversation. Enter utterances that you expect to trigger your action, and then adjust the copilot action and topic instructions based on your results. Your action must be assigned to your copilot to test it.

Best Practices for Einstein Copilot Action Instructions

Action instructions tell a copilot what your action does and when to use it in a conversation. When you create a custom action, carefully plan and test your instructions to ensure that your action performs as expected.

REQUIRED EDITIONS

Available in: Lightning Experience

Available in: Enterprise, Performance, and Unlimited Editions with the Einstein for Sales, Einstein for Service, or Einstein Platform add-ons.

The names and instructions for copilot actions and their inputs and outputs determine when and how an action is used in a conversation. Einstein Copilot uses a large language model (LLM) and the planner service to search through all actions assigned to a copilot and select one or more to execute in the conversation, based on the user's question or request. The more effective the instructions, the more predictably the action performs.

Effective instructions vary by action and use case. For best results, test your action in the Agent Builder and iterate until your action meets your requirements.

Instructions for Actions

Action instructions describe what your action does and when to use it in a conversation. The most important element of action instructions is a concise natural language description of what your action does, usually one to three sentences written as an instruction to the LLM. The description can include the goal or task of the action, any use cases for the action, and the objects or records it uses or modifies. In general, the more relevant detail you include in your instructions, the easier it is for the LLM to differentiate between actions.

EXAMPLE COPILOT ACTION	LESS EFFECTIVE DESCRIPTION	MORE EFFECTIVE DESCRIPTION
Update Contact Phone Number	Updates a phone number.	Updates the user's phone number associated with their contact record. If no matching contact exists, creates a new contact record.
Query Records	Provides the ability to use natural language to query information about Salesforce CRM records.	Finds and retrieves Salesforce CRM records based on user input and specific conditions, such as field values. This action automatically identifies the correct records and object type.
Track Order	Action to get tracking information for an order.	Returns tracking information for a customer order based on the tracking number and destination ZIP code.

 TIP LLMs work best when you use unique or varied language. When naming copilot actions, instead of starting similar actions with the same verb (for example, "Get"), try a variety of related verbs (for example, "Get," "Find," "Identify," "Retrieve"). When writing copilot action instructions, remove or revise nonessential phrases that are repeated across multiple actions (for example, "This code snippet....," "Provides the ability to...").

In many cases, the natural language description is all that's required for a copilot to use your action as expected. However, some actions require more details after testing. Here are some options for refining your action description.

WHAT TO INCLUDE	WHEN TO USE	EXAMPLES
One to three utterances, or example customer inputs	Include examples of customer messages that you expect to trigger your action to help the LLM match a given customer input to your action. If you include multiple utterances, minimize overlap. <ul style="list-style-type: none">Vary the length, grammar, and word choice of your utterances.If it makes sense for your use case, phrase one utterance as a statement and another as a question.Provide different utterances for the different ways your action can be used (for example, updating fields on a case versus an account or drafting versus revising an email).	<ul style="list-style-type: none">"What's the address for the Acme account?""Show me all open opportunities with a close date before 10/31/24.""What's your return policy?""How do I fix an issue?""What features does a product have?"
Dependent actions	Some actions use other actions, often system actions, to get required information. If a required action is missing from a plan or actions aren't being called in the right order, you can specify them by API name.	The action <code>IdentifyRecordByName</code> must be called directly prior to this action to get a single record ID as input. If the <code>IdentifyRecordByName</code> action is part of the plan, it must run before <code>GetRecordFieldsAndValues</code> .
Examples of how you want the action to function	Examples don't generalize well, but they can be helpful when a copilot makes a consistent mistake. In general, LLMs understand examples of what to do better than examples of what not to do, but you can test with either.	This action should be called only when a more specific summarization action isn't available.

Instructions for Inputs and Outputs

Instructions for inputs and outputs describe how to format and use the parameter. In the instructions for an input, specify what information is required to complete the action. In the instructions for an output, specify the result of the action and what a copilot returns to the user.

LLMs don't have a complete understanding of context, including the Salesforce context. In other words, an LLM doesn't always know the number of digits in a Salesforce record ID or the names of Salesforce objects. Concise but detailed instructions help a copilot retrieve necessary information and return relevant results. The best way to determine the right level of detail is to test and iterate on your instructions.

When writing instructions for an input or an output, consider:

- The data type of the information: "an 18-digit ID," "a rich-text answer to the user's question."
- Any specific fields or values: "contains the contact ID and expected delivery date".
- How the information is retrieved or generated: "For a new email, retrieve the ID from the user's input. For a revised email, retrieve the ID from the conversation history."
- Any special rules about how you expect the action to be used: "Don't generate SQL or SOQL", "Required only when the user wants to revise an earlier draft.", "The API name is used to identify Salesforce record field-value pairs from the user request."

Here are some examples of effective instructions for inputs and outputs.

INPUT OR OUTPUT PARAMETER	INSTRUCTIONS
recordId	The single ID of the Salesforce CRM record to create the summary for. For example: recordId, record_id, recordIds[0], accountId.
objectApiName	The API name of the Salesforce object (such as Account or Opportunity) associated with the record the user wants to find. The name can be obtained from the context and doesn't require the IdentifyObjectByName action.
summary	The rich text summary created for the specified record.
searchResults	A list of the matching Salesforce CRM record IDs in descending order of relevance.
generatedEmail	A JSON object that contains the recipient ID, subject line, and body of the email.

Explore Standard Copilot Topics and Actions

Get to know Salesforce standard copilot topics and actions with our comprehensive reference libraries, so you can quickly add powerful out-of-the-box functionality to your copilot.

REQUIRED EDITIONS

Available in: Lightning Experience

Available in: Enterprise, Performance, and Unlimited Editions with the Einstein for Sales, Einstein for Service, or Einstein Platform add-ons.

- Standard Topic Reference**

Standard topics are topics available out of the box in the asset library. Many standard topics are available to all users who have permission to access Einstein Copilot. Some standard topics are built to work with specific clouds or licenses, so they require another license or permission.

- Standard Copilot Action Reference**

Standard actions are included by default with your copilot. Many standard copilot actions are available to all users who have permission to access Einstein Copilot. Other standard actions are built to work with specific clouds or products, so they require an additional license or permission.

Standard Copilot Topics

Explore the comprehensive reference library for Einstein Copilot standard topics.

- Copilot Topic: Close Deals**

The Close Deals topic leverages AI-powered insights to heighten sales success. Get recommendations on the best approach to close deals, based on past deals won, conversation signals, and customer sentiment. Sellers also get answers to product-related pricing questions and suggestions for close plans to help you close deals.

- Copilot Topic: Communicate with Customers**

Generate personalized and clear communications to help sellers engage with prospects, contacts, and leads via email, follow-ups, and meeting requests.

- Copilot Topic: Conversation Explorer**

Answer user's questions about a voice or video call based on the contents of the call transcript.

- Copilot Topic: Field Service Dispatcher Actions**

This topic helps users find appointments that require immediate attention and shows them in the appointment list for subsequent follow-up actions.

- Copilot Topic: Forecast Sales Revenue**

Use the Forecast Sales Revenue topic for comprehensive insight into predicting sales revenue. Gain an understanding of your team's forecasts, predict future sales based on current market signals or predictive analysis, and get answers to queries related to deal alerts, risky deals, or estimations around deal value.

- Copilot Topic: General CRM**

Handles user requests that are related to Salesforce CRM data, such as identifying, summarizing, or updating records, drafting or refining emails, aggregating data, and finding and querying Salesforce objects.

- Copilot Topic: Manage Deals**

The Manage Deals topic equips sellers with the strategic information to effectively oversee your deals. Sellers can get information about specific contacts or relationships with contact accounts, empowering them to make informed decisions. By identifying similar opportunities, users can optimize their deal strategies and prioritize crucial activities and accounts, all while ensuring their CRM remains current with the latest tasks, calls, and customer interactions. This comprehensive approach provides a more real-time view of deals and pipelines, ultimately leading to more effective deal management.

- Copilot Topic: Single Record Summary**

Generates a summary of a single record.

Standard Copilot Actions ↴

Explore the comprehensive reference library for Einstein Copilot standard actions.

- **Copilot Action: Add Record to Cadence**
Associates a prospect with a cadence directly from Einstein Copilot.
- **Copilot Action: Answer Questions with Knowledge**
Answers a question from a user based on information from relevant knowledge articles.
- **Copilot Action: Create a Label**
Creates a personalized label to contextually group and organize records.
- **Copilot Action: Create a To-Do**
Create tasks for follow-up or assignments to complete, helping sellers manage their work efficiently.
- **Copilot Action: Create Close Plan**
Generates a sales plan to help sales reps close an opportunity by its target date.
- **Copilot Action: Create Appointment List Filter**
Creates a filter for the service appointment list in the dispatcher console. If a user triggers QueryRecords to retrieve service appointments with certain criteria and then asks to see the service appointments in the appointment list, the action creates a filter in the Appointment List containing the service appointments returned by Copilot. This action is one of the actions provided by the Einstein Copilot: Field Service Actions feature
- **Copilot Action: Draft or Revise Email**
Creates or revises a draft of an email to a recipient based on the user's request and available email types. Email drafts include a recipient, subject, and email body. Email drafts can be sent through the Salesforce email composer or copied to another email client.
- **Copilot Action: Explore Conversation**
Answers a user's questions about a voice or video call based on the contents of the call transcript. For example, if the user enters, "What was the customer sentiment on the call?" in the Einstein panel, the action delivers an answer based on the customer's language in the call transcript.
- **Copilot Action: Find Contact Interactions**
Identifies the list of opportunities or an account that a specific contact has influenced.
- **Copilot Action: Find Similar Opportunities**
Searches for and returns a list of opportunities that are similar to, but not duplicates of, a specified Salesforce opportunity. An explanation of the similarities is returned with each opportunity record.
- **Copilot Action: Get Activities Timeline**
Retrieves a list of all CRM activities associated with a record during the specified time frame. The list can include past and future activities.
- **Copilot Action: Get Activity Details**
Provides a brief summary of an activity, including an overview of its content (such as an email or a call summary) and other relevant details. Valid activity types are calls, emails, events, or tasks.
- **Copilot Action: Get Forecast Guidance**
Gathers information about how a sales rep is tracking toward meeting their Committed and Most Likely quota for the current forecast period. The output includes the seller's forecast amount, a list of any deals at risk, and why they're considered at risk.
- **Copilot Action: Get Product Pricing**
Gets the pricing information of a product, including relevant historical sale price data from previous won deals involving the same product.
- **Copilot Action: Get Record Details**
Retrieves the details of a record, including object fields and values, records from related lists, and associated tasks and events.
- **Copilot Action: Identify Contact Role**
Identifies the role, significance, or impact of a contact in influencing the opportunity or account status. This action also assesses how that contact affects the outcome of a deal, initiative, or account.
- **Copilot Action: Identify Key Contacts**
Identifies key contacts associated with an account or opportunity. A key contact is someone who can influence, obstruct, impede, or decide a deal, or is involved in the outcome of the deal.
- **Copilot Action: Identify Object by Name**
Interprets the user's input to determine which object the user is referring to and then returns a list of matching object names so other actions can be taken. For example, if the user enters, "List the opportunities for the Acme account" in the Einstein panel, the action identifies that the user is requesting information related to the Account object and Opportunity object.
- **Copilot Action: Identify Record by Name**
Searches for Salesforce records by name and returns a list of matching record IDs. For example, if a user enters, "Show me the Acme records" in a Salesforce org with multiple records related to "Acme," the action can return the Acme account, the Acme - 20k opportunity, and the Acme - 50k opportunity.
- **Copilot Action: Label a Record**
Assigns personalized labels to records. Users can tag records based on meaningful attributes, retrieve those records as a group, and act on them.
- **Copilot Action: Log a Call**
Logs a call with a customer as a task, which can include additional details about the call, such as the date, time, and purpose. The logged call appears in the seller's Activity Timeline.
- **Copilot Action: Prioritize Opportunities**
Identifies a rep's top sales opportunities and ranks them in order of priority based on specific instructions provided by the user or available opportunity data. An explanation for each opportunity's rank is also returned.
- **Copilot Action: Query Records (Beta)**
Finds and retrieves Salesforce records based on the user's request and specific conditions, such as the values of fields. The action also generates an explanation describing how the user's request was interpreted to retrieve the search results.
- **Copilot Action: Query Records with Aggregate (Beta)**
Answers aggregation questions about Salesforce data, such as count, sum, average, max, min, and group-by.
- **Copilot Action: Send Meeting Request**
Creates an email draft to request a meeting with the recipient, with three possible time slots, a link to view the slots on a new page, or a link to the sender's calendar where they can pick a slot to schedule a meeting.
- **Copilot Action: Summarize Record**
Generates a summary of a record. The output depends on the record type and the prompt templates the user has access to.
- **Copilot Action: Summarize Service Appointment Notes**
Generates a summary of one or more service appointment notes. The output depends on the prompt templates the user has access to.
- **Copilot Action: Summarize Scheduling Issues for Field Service**
Creates a summary of the total number of appointments that require immediate attention. The summary is divided according to exception categories. The copilot action creates a filter for each category in the appointment list. This action is one of the actions provided by the Einstein Copilot: Field Service Actions feature. The categories used by the reports are customizable.
- **Marketing Cloud Copilot Actions**
Use these actions to create briefs, campaigns, and message content in Marketing Cloud.

Einstein Data Library

Improve accuracy, add personalization, and build trust in Einstein's responses when you add data libraries to your Einstein features. Einstein Data Library uses grounding to index your Knowledge articles and fields so that Einstein knows which information to base responses on. Using this index, Einstein can quickly check the accuracy of responses against your Knowledge base, so you get the best results.

REQUIRED EDITIONS

Available in: Lightning Experience

Available in: Enterprise and Unlimited Editions with the Einstein for Sales, Einstein for Platform, or Einstein for Service add-on

- **Setting Up Data Libraries**

Use Einstein Data Library to enhance the accuracy of AI in your features.

Einstein in Sales and Service Cloud

* How to increase sales reps productivity?

- **Einstein Activity Capture for Sales Cloud Einstein**

Let Einstein help keep data between Salesforce and your email and calendar applications up to date. Then you can get back to more important things, like selling.

- **Einstein Automated Contacts**

Spend even less time on data entry. Einstein Automated Contacts uses email and event activity to find new contacts and opportunity contact roles to add to Salesforce. Depending on how your Salesforce admin set up the feature, Einstein suggests the new data or adds it for you.

- **Einstein Generative AI for Sales**

Einstein Generative AI for Sales connects your sales data with the generative capabilities of artificial intelligence (AI) and large language models (LLMs). With Sales Emails, your sales team can generate personalized emails to send to contacts and leads. Your team gets precious time back in their day by getting nearly-there drafts, already contextualized and grounded in your Salesforce data. And because trust is our #1 value, you can rest easy knowing that your data is secure.

* **Call Summary**: Powered by Einstein Gen AI. User can generate summary of a audio and video calls. The new Summary tab allows user to create editable summary that includes next steps and customer feedback.

* **Call Explorer**: User can ask questions directly from voice and video call records.

* **Opportunity Scoring**: Provides insights on oppy instead of Leads.

* **Einstein Activity Capture**: Automates the logging of sales activities into Salesforce.

* **Einstein Lead Scoring**: Scores leads by how well they fit into company's successful conversion pattern.

* **Einstein Service Replies** → Feature of Service Cloud that generates email and chat responses based on knowledge base data.

* **Einstein Automated Contacts**: Finds new contacts and oppy contact roles to add to Salesforce.

* **Einstein Case Classification**: Automates the categorization of case.

* **Einstein Service Replies for Email**: Helps user to draft and send personalized email response to customers based on recommended knowledge articles.

* **Einstein Service Replies for Chat**: Draft replies to customers during live chat or messaging sessions.

- * **Einstein Work Summary**: Provides real time summaries of ongoing conversation including the issue and resolution.
- * **Einstein Article Recommendations**: Recommends knowledge articles that were attached to similar cases in the past.

Model Builder: Using model builder you can create your own predictive model with clicks.

Predictive Models address common business use cases

1. **Regression for Numbers** → Numeric use cases which targets number/numeric outcome represented as quantitative data such as currency, counts or percentages.

Examples - Predict Amount of an Opppty.

- Time to close —
- customer lifetime value of an Account —
- satisfaction of a case —

2. **Binary Classification for Text**: Binary classification solution target business outcomes with only two possible results which are represented as text data.

Outcomes are typically yes/no questions that are expressed in business terms such as churned or not churned, Opppty won or lost, employee retained or not retained. For analysis purpose Einstein will convert values to true & false.

Examples : Predict the possibility to win an Opppty.

- of an account to buy product. —
- that a lead will convert. —
- that an Account will churn. —

Use Model Builder to create a model that uses historical data to predict future outcomes. Select and structure your data source, define your desired outcome, and train your model to reveal predictive inferences and insights using AI, machine learning, and statistical analysis.

The data source must meet these requirements.

REQUIREMENT	MINIMUM	MAXIMUM
Number of rows	400	20 million 5 million for the XGBoost algorithm
Number of columns	3 (1 outcome variable plus 2 other columns)	50

